



FORCES NLP: an efficient implementation of interior-point methods for multistage nonlinear nonconvex programs

A. Zanelli, A. Domahidi, J. Jerez & M. Morari

To cite this article: A. Zanelli, A. Domahidi, J. Jerez & M. Morari (2020) FORCES NLP: an efficient implementation of interior-point methods for multistage nonlinear nonconvex programs, International Journal of Control, 93:1, 13-29, DOI: [10.1080/00207179.2017.1316017](https://doi.org/10.1080/00207179.2017.1316017)

To link to this article: <https://doi.org/10.1080/00207179.2017.1316017>



Published online: 22 May 2017.



Submit your article to this journal [↗](#)



Article views: 1735



View related articles [↗](#)



View Crossmark data [↗](#)



Citing articles: 20 View citing articles [↗](#)



FORCES NLP: an efficient implementation of interior-point methods for multistage nonlinear nonconvex programs

A. Zanelli^a, A. Domahidi^{b,c}, J. Jerez^{b,d} and M. Morari^d

^aSystemtheorie, Regelungstechnik und Optimierung, Albert-Ludwigs-Universität, Freiburg, Germany; ^bembotech GmbH, Zurich, Switzerland; ^cinspire AG, Zurich, Switzerland; ^dAutomatic Control Laboratory, ETH Zurich, Zurich, Switzerland

ABSTRACT

Real-time implementation of optimisation-based control and trajectory planning can be very challenging for nonlinear systems. As a result, if an implementation based on a fixed linearisation is not suitable, the nonlinear problems are typically locally approximated online, in order to leverage the speed and robustness of embedded solvers for convex quadratic programs (QP) developed during the last decade. The purpose of this paper is to demonstrate that, using simple standard building blocks from nonlinear programming, combined with a structure-exploiting linear system solver, it is possible to achieve computation times in the range typical of solvers for QPs, while retaining nonlinearities and solving the nonlinear programs (NLP) to local optimality. The implemented algorithm is an interior-point method with approximate Hessians and adaptive barrier rules, and is provided as an extension to the C code generator FORCES. Three detailed examples are provided that illustrate a significant improvement in control performance when solving NLPs, with computation times that are comparable with those achieved by fast approximate schemes and up to an order of magnitude faster than the state-of-the-art interior-point solver IPOPT.

ARTICLE HISTORY

Received 30 June 2016
Accepted 2 April 2017

KEYWORDS

Predictive control; nonlinear systems; numerical optimisation; interior-point methods; embedded systems

1. Introduction

Model predictive control (MPC) has drawn a lot of attention in both academia and industry since the late 1970s (Garcia, Prett, & Morari, 1988). The possibility of formulating an optimisation problem that directly encodes control objectives and constraints, together with its inherent multivariable nature, has made MPC a promising technology for the systematic design of control systems in several fields (Qin & Badgwell, 2003). This has become possible thanks to both the availability of reliable and efficient numerical methods and the increasing computational capabilities of embedded computing units. Despite the increased complexity of the algorithms in comparison with classical control techniques, considerable progress has been made in the last decade to make MPC a viable technology in many challenging applications. For linear systems, fast implementations of convex solvers running at millisecond and microsecond timescale (Domahidi et al., 2012; Ferreau, Kirches, Potschka, Bock, & Diehl, 2014; Frison, Sorensen, Dammann, & Jørgensen, 2014; Jerez et al., 2014; Richter, Jones, & Morari, 2012) have been developed and, for some methods, complexity certifications have been derived (Giselsson, 2012; Richter et al., 2012), providing bounds on the maximum number

of iterations to guarantee convergence to a global optimum.

As a result, methods for the solution of convex quadratic programs (QPs) have become a reliable technology. When handling nonlinear nonconvex problems, it is common practice to approximate nonlinearities, potentially present in the system dynamics, constraints and objectives, in order to rely on existing methods and implementations for linear-quadratic MPC. A common approach to handle nonlinear models is to linearise the dynamics around the current approximation of the optimal trajectories. This procedure provides a linear time-varying model that locally approximates the dynamics and can be readily used to formulate a linear time-varying MPC problem as a convex QP. This strategy is employed, for example, in Falcone et al. (2007), which addresses the problem of actively controlling the front steering system of an autonomous vehicle. Alternatively, in Kothare, Balakrishnan, and Morari (1996) the effect of the nonlinear dynamics is taken into account by formulating a robust optimisation problem. The original model is approximated by a polytopic uncertain linear time-varying system. Prior information on the Jacobian of the dynamics is used to guarantee that the trajectories of the system are confined inside polytopes. In this way a

min-max problem is formulated that can be solved with semidefinite programming methods. However, the resulting scheme is potentially highly conservative and still too computationally intensive for most practical purposes. Bemporad and Morari (1999) propose instead to approximate the dynamics of a nonlinear system with a piecewise affine function, requiring the introduction of integer variables into the optimisation problem, which greatly increases the computational complexity of the solution algorithms. Alternative approaches based on feedback linearisation are proposed, for example, in Simon, Lofberg, and Glad (2013) and Deng, Becerra, and Stobart (2009). Feedback linearisation results in problems with nonlinear and, in general, nonconvex inequality constraints that require specific handling. Despite the satisfactory results reported in some specific cases, these approaches, in general, result in formulations that are computationally infeasible for real-time applications or degrade the control performance significantly.

Including nonlinearity directly into the optimal control problem allows one to circumvent these difficulties and provides a systematic approach for handling systems with nonlinear dynamics, constraints and objectives. Following this path relies on the ability to efficiently and reliably solve the nonlinear and nonconvex optimisation problems arising from an optimal control formulation. While there exist approaches in the literature that rely on dynamic programming techniques or indirect methods (Diehl, Ferreau, & Haverbeke, 2009), most existing strategies are nowadays based on the so-called *direct methods* (Diehl et al., 2009), which formulate a finite-dimensional discrete-time nonlinear program (NLP). State-of-the-art numerical solvers for nonlinear programming can be used to solve the resulting discretised formulations.

Compared to the convex QPs arising in linear-quadratic MPC, optimisation problems arising from nonlinear MPC formulations present several additional challenges that can have a major impact on both reliability and efficiency of solvers. This is especially important for embedded systems that have limited computational power and must run autonomously without user interaction. First, the discretisation of the continuous-time problem requires the solution of nonlinear algebraic differential equations and the computation of derivatives for the current predicted trajectories with respect to the optimisation variables, two operations that can be computationally demanding. Second, the discretised problem is in general nonconvex and a number of additional factors must be taken into account in order to guarantee convergence of the solver to a local minimum. This leads to a significant increase in the complexity of the methods.

Due to the high computational burden associated with the solution of nonlinear optimisation problems, several approximate schemes have been proposed to trade control performance for speed (Diehl et al., 2009). Among these schemes, the *real-time iteration* (RTI) proposed in Diehl et al. (2002) is based on the idea of refining the solution while the problem changes (Diehl et al., 2009). A single convex QP that locally approximates the original optimisation problem is solved per feedback step. If the iterates are initialised sufficiently close to a local minimum, a solution is gradually approached, while controlling the system using the current approximation of the solution. Local stability properties of such a scheme have been investigated in Diehl, Findeisen, and Allgöwer (2007). A region of attraction can be defined such that, if the combined system-optimiser state lies in it, the system can be stabilised using the mentioned approximate feedback law. When using this method, stability problems can arise if, after a large disturbance, the system leaves the region of attraction (Houska, Ferreau, & Diehl, 2011). A second approximate scheme in the literature is the so-called *Continuation-GMRES* (C-GMRES) method in Ohtsuka (2004). As for the RTI, closed-loop stability of C-GMRES is in principle covered by the considerations in Diehl et al. (2007) and analogous issues can be encountered in the presence of large disturbances or reference changes.

In order to avoid the eventual stability and suboptimality issues related to approximate schemes, it is possible to solve the optimisation problems to a local minimum. In this way, classical results on nominal and robust stability (Nicolao, Magni, & Scattolini, 1998) can be applied. However, due to long computation times, feedback delays can degrade control performance in practice. Practical and theoretical issues have been discussed in the literature (Chen, Balance, & O'Reilly, 2000; Findeisen & Allgöwer, 2004; Santos, Afonso, Castro, Oliveira, & Biegler, 2001). Zavala and Biegler (2009) propose the so-called *advanced-step NMPC* controller to overcome these difficulties. The main idea is to predict the state of the system in order to solve the future optimal control problem in advance, hence compensating for the computational delay. Before applying the computed control input to the system, a linearisation of the solution manifold is used to perform a computationally cheap correction step based on the actual state estimate or measurement. This approach has been proven to perform well if a good prediction of the state of the system is available, while a considerable performance degradation is experienced in case of plant-model mismatch or large disturbances (Zavala and Biegler, 2009).

Main contribution and outline

Ideally, one would not need to trade-off control performance for computational delay. This paper presents a software package that implements an efficient interior-point method for nonlinear and nonconvex optimisation problems arising from optimal control formulations. The main contributions are the following:

- (1) The main features of the C code generator FORCES NLP (Domahidi & Jerez, 2016) are presented. The multistage structure of the NLPs is exploited in order to efficiently build Hessian approximations and solve the Karush–Kuhn–Tucker (KKT) systems to compute search directions. In particular, a block-wise Cholesky factorisation is used based on the approach in Domahidi et al. (2012). The class of problems solvable by the structure-exploiting QP solver in Domahidi et al. (2012) is extended to the more general class of smooth nonlinear nonconvex programs.
- (2) Most available methods and software packages capable of achieving fast feedback times rely on approximate solution schemes. Despite the good control performance attainable in many cases, the system is often operated suboptimally. The proposed approach solves each optimisation problem to a local minimum exploiting an efficient interior-point method in order to reduce the computational burden. The numerical results in Section 4 show how the resulting control performance can be considerably improved, especially for strongly nonlinear problems, with moderately higher computation times.
- (3) Thanks to the structure-exploiting linear algebra, a considerable speedup can be achieved with respect to state-of-the-art solvers for nonlinear programming. A comparison with the state-of-the-art code IPOPT shows the large computational benefits of the proposed implementation for multistage problems.

The remainder of this paper is organised as follows: in Section 2, a brief introduction to nonlinear MPC is given and state-of-the-art numerical methods for the solution of the arising optimisation problems are outlined. Section 3 discusses the implementation details of the software package presented and in Section 4 numerical results are discussed that compare the solver with two different schemes: the fast and approximate RTI scheme and the exact, but generally computationally demanding, interior-point solver IPOPT (Wächter & Biegler, 2006).

2. Preliminaries

2.1 Nonlinear model predictive control

2.1.1 Problem formulation

In this work the following continuous-time, infinite-dimensional optimisation problem will be considered:

$$\begin{aligned} \min_{x(\cdot), u(\cdot)} \quad & \int_{t_0}^{t_f} l_c(t, x(t), u(t)) dt + e(x(t_f)) \\ \text{s.t.} \quad & x(t_0) = x_0 \\ & \dot{x}(t) = f(t, x(t), u(t)) \\ & g(t, x(t), u(t)) \leq 0, \end{aligned} \quad (1)$$

where $x \in \mathbb{R}^{n_x}$ are the states of the system, $u \in \mathbb{R}^{n_u}$ are its inputs and l_c , e , f and g are potentially nonlinear twice continuously differentiable functions. In order to recast the problem into a finite-dimensional program that can be solved within practical computation times for large system dimensions, the following reformulations are applied:

- (1) As closed-form solutions for nonlinear differential equations are generally not available, the state trajectories need to be approximated numerically using an integration scheme;
- (2) A piece-wise constant parametrisation of the control trajectories is used that results in a problem with a finite number of variables.

This so-called direct approach results in a formulation that can be solved with state-of-the-art methods for nonlinear optimisation that can handle large-scale problems and can deal with constraints directly, avoiding the main difficulties related to indirect methods and approaches based on dynamic programming. Notice that these problems are in general nonconvex, hence computing the global minimum with direct methods can be a computationally intensive task. Most global optimisation algorithms require the solution of several local NLPs (Nocedal & Wright, 2006), which is often not practical for embedded applications.

Among several approaches available to discretise problem (1), *direct multiple shooting* (Bock & Plitt, 1984) will be considered in this work, due to the special structure present in the obtained discretised problems. State trajectories at time $t = t_0, \dots, t_N$ are approximated using a numerical integration scheme, obtaining the following discrete-time finite-dimensional problem,

$$\begin{aligned} \min_{z_0, z_1, \dots, z_N} \quad & \sum_{i=0}^N l_i(z_i) \\ \text{s.t.} \quad & c_i(z_{i+1}, z_i) = 0, \quad i = 0, \dots, N-1 \\ & h_i(z_i) \leq 0, \quad i = 0, \dots, N, \end{aligned} \quad (2)$$

with stage variables $z_i \in \mathbb{R}^{p_i}$, stage cost functions $l_i : \mathbb{R}^{p_i} \rightarrow \mathbb{R}$, and equality constraints $c_i : \mathbb{R}^{p_i+p_{i+1}} \rightarrow \mathbb{R}^{r_i}$. General twice differentiable functions l_i and c_i will be considered and only simple bounds

$$h_i := \begin{bmatrix} \bar{z}_i - z_i \\ z_i - \underline{\bar{z}}_i \end{bmatrix} \quad (3)$$

on the stage variables are present in the discretised formulation.

Notice that problems with general nonlinear inequalities can be recast into form (2) by introducing additional variables and equality constraints. Hence, although problems with general inequalities can be directly specified from the interface of FORCES NLP, in the following, formulation (2) will be considered, without loss of generality. The structure of (2) is the same as in convex QPs arising from linear-quadratic MPC formulations and it can be exploited in order to efficiently compute the Newton search directions in an interior-point method, as outlined in Section 3.

The continuous-time formulation (1) can be discretised using other schemes as well. The so-called *direct single shooting* scheme provides a more compact formulation obtained by simulating the dynamics in order to eliminate the state variables. The resulting problem has fewer variables than the one obtained by applying multiple shooting. On the other hand, the favorable multistage sparsity structure is lost (Diehl et al., 2009). A different approach consists in formulating a problem that embeds the nonlinear equations used to discretise the dynamics, for instance, the equations of an implicit Runge-Kutta integration scheme including all intermediate variables. Such a scheme, referred to as *direct collocation*, results in a larger optimisation problem with an alternative structure that can be exploited using different approaches to the ones used in this paper (Kang, Cao, Word, & Laird, 2014; Quirynen, Gros, Houska, & Diehl, 2016). Further details on integration schemes used to obtain formulation (2) from the continuous-time optimal control problem (1) are given in Section 3.5.

2.2 Newton-type approaches to nonlinear programming

Newton-type approaches to nonlinear programming can be interpreted as variations of the Newton method applied to the KKT system associated with (2),

$$r_S := \nabla_z \mathcal{L} = 0 \quad (4a)$$

$$r_C := c(z) = 0 \quad (4b)$$

$$r_D := h(z) + s = 0 \quad (4c)$$

$$r_N := Y_D s = 0 \quad (4d)$$

$$y_d \geq 0, \quad s \geq 0, \quad (4e)$$

where

$$\mathcal{L} := \sum_{i=0}^N l_i(z_i) + y_c^T c(z) + y_d^T h(z) \quad (5)$$

is the Lagrangian of the problem, y_c and y_d are the multipliers associated with equality and inequality constraints, respectively, s are slack variables and Y_D is the diagonal matrix having the elements of y_d on its diagonal.

2.2.1 Sequential quadratic programming

A first approach that can be used in order to numerically solve problem (2) is sequential quadratic programming (SQP). The main idea consists in iteratively approximating the nonlinear problem with QPs. Given a candidate solution z^k the local quadratic approximation takes the form

$$\begin{aligned} \min_{\Delta z_0, \dots, \Delta z_N} \quad & \sum_{i=0}^N \left(\frac{1}{2} \Delta z_i^T H_i \Delta z_i + \nabla l_i(z_i^k)^T \Delta z_i \right) \\ \text{s.t.} \quad & c_i(z_{i+1}^k, z_i^k) + \nabla c_i(z_{i+1}^k, z_i^k)^T [\Delta z_i^T \Delta z_{i+1}^T]^T = 0, \\ & i = 0, \dots, N-1, \\ & h_i(z_i^k) + \nabla h_i(z_i^k)^T \Delta z_i \leq 0, \\ & i = 0, \dots, N, \end{aligned} \quad (6)$$

where H_i is an approximation of the block of the Hessian of the Lagrangian associated with z_i . After solving (6), the solution is updated

$$z^{k+1} = z^k + \Delta z \quad (7)$$

and it can be shown that, under mild assumptions, the iterates converge to a local solution of the discretised optimisation problem (Nocedal & Wright, 2006).

Due to the potentially high computational burden associated with the solution of several QPs and the computation of several linearisations, there exist approaches in the literature that rely on a limited number of iterations. The RTI scheme is based on this idea. It only performs a single step of an SQP-based method at each sampling instant and the resulting approximate solution is used to control the system. Although performance degradation can be experienced due to the suboptimality of the approximate solution, this approach has been proven to work well in several practical cases (Ferreau, Houska, Geebelen, & Diehl, 2011; Quirynen, Gros, & Diehl, 2013a).

2.2.2 Interior-point methods

An alternative approach consists of applying the Newton method directly to the set of nonlinear equations (4). In order to do so, the nonsmooth condition (4d) is relaxed, such that a set of smooth nonlinear equations is obtained to which the Newton method can be applied directly. Once a solution to the relaxed KKT system is obtained, the relaxation parameter τ is decreased and, as $\tau \rightarrow 0$, a local optimum is approached (Nocedal & Wright, 2006). The main computational cost is associated with evaluating functions and their derivatives, and computing search directions. At every iteration of the interior-point algorithm the KKT system is linearised at the current iterate and a linear system is solved to obtain a search direction. For stiff differential equations and differential algebraic equations, a high computational burden can be associated with function evaluations and derivatives generation. Recent results on the efficient implementation of numerical methods for the integration of such problems (Quirynen, Vukov, & Diehl, 2012; Quirynen, Gros, & Diehl, 2013a, 2013b) can be exploited in order to considerably reduce computation times required for linearisation. In the following section the problem of efficiently computing the search directions is addressed.

3. Implementation

3.1 Structure-exploiting linear system solver

An efficient interior-point method is used to solve (4) based on the implementation proposed in Domahidi et al. (2012). While the KKT solver in Domahidi et al. (2012) targets convex QPs and convex quadratically constrained quadratic problems (QCQPs), in this work an extension that can handle general nonlinear and nonconvex multi-stage problems is presented.

The linearised KKT version of (4) for problem (2) is

$$\begin{bmatrix} \mathcal{H}(z, y_c) & J_{eq}(z)^T & J_{ineq}(z)^T \\ J_{eq}(z) & & \\ J_{ineq}(z) & & \end{bmatrix} \begin{bmatrix} \Delta z \\ \Delta y_c \\ \Delta y_d \\ \Delta s \end{bmatrix} = - \begin{bmatrix} r_s \\ r_c \\ r_d \\ r_N \end{bmatrix}, \quad (8)$$

where $\mathcal{H}(z, y_c)$ is block-diagonal with blocks given by

$$\begin{aligned} \mathcal{H}_i &:= \nabla_{z_i}^2 l_i(z_i) + \sum_{k=1}^{r_{i-1}} y_{c,i-1[k]} \nabla_{z_i}^2 c_{i-1[k]}(z_i, z_{i-1}) \\ &+ \sum_{k=1}^{r_i} y_{c,i[k]} \nabla_{z_i}^2 c_{i[k]}(z_{i+1}, z_i), \end{aligned}$$

where r_i and q_i are the number of equality and inequality constraints in stage i , respectively, and k denoting the k th

entry or constraint. J_{ineq} is also block diagonal with blocks consisting of $\nabla h_i(z_i)$, while J_{eq} is block banded with rows consisting of $\nabla c_i(z_{i+1}, z_i)$ in the appropriate position. Moreover, $S \in \mathbb{R}^{q \times q}$ has been introduced to denote the diagonal matrix having s on its diagonal.

When dealing with nonlinear problems instead of convex QPs, \mathcal{H} is a function of the primal variables z , but also of the equality multipliers y_c , due to the potential nonlinearity in the dynamics and constraints. The Jacobians of the equality and inequality constraints, for the same reason, are functions of z rather than being constant as in QPs and QCQPs. Moreover, it is important to mention that it is in general difficult to ensure that the KKT matrix in (8) has the desired inertia, i.e. the right number of positive, negative and zero eigenvalues, which is required to guarantee important descent properties (Wächter & Biegler, 2006). With the Hessian of the Lagrangian potentially being indefinite, additional care has to be taken when solving system (8) as outlined in the following subsection.

The search direction (8) is computed with a direct structure-exploiting factorisation method. In particular, the linear system is first reduced by eliminating Δs :

$$\Delta s = Y_D^{-1} (-r_N - S \Delta y_d) \quad (9)$$

and successively Δy_d :

$$\Delta y_d = S^{-1} Y_D (r_D + J_{ineq}(z) \Delta z) - S^{-1} r_N, \quad (10)$$

obtaining the following symmetric indefinite system:

$$\begin{bmatrix} \Phi & J_{eq}(z)^T \\ J_{eq}(z) & 0 \end{bmatrix} \begin{bmatrix} \Delta z \\ \Delta y_c \end{bmatrix} = - \begin{bmatrix} r_d \\ r_c \end{bmatrix}, \quad (11)$$

where

$$\begin{aligned} \Phi &:= \mathcal{H}(z, y_c, y_d) + J_{ineq}(z)^T S^{-1} Y_D J_{ineq}(z) \\ r_d &:= r_s + J_{ineq}(z)^T S^{-1} Y_D r_D - J_{ineq}(z)^T S^{-1} r_N. \end{aligned} \quad (12)$$

This is the form of the KKT system used, for example, in Wächter and Biegler (2006), which will be referred to as the *symmetric indefinite* form. As done in Domahidi et al. (2012), a more compact and symmetric system can be obtained by forming the Schur complement of the iteration matrix (11):

$$Y \Delta y_c = \beta, \quad (13)$$

where

$$\begin{aligned} Y &:= J_{\text{eq}}(z) \Phi^{-1} J_{\text{eq}}(z)^T \\ \beta &:= r_C - J_{\text{eq}}(z) \Phi^{-1} r_d \\ \Delta z &= \Phi^{-1} (-r_d - J_{\text{eq}}(z)^T \Delta y_c). \end{aligned} \quad (14)$$

This form of the KKT system is also referred to as the *normal equations* form.

Another additional complication with respect to QPs arises because the Hessian and Jacobian matrices forming the linearised KKT matrix (8) are functions of the current iterate. As a result, matrices Φ and $J_{\text{eq}}(z)$ can become rank deficient. In general, it is necessary to add regularisation terms to the Hessian and equality constraint Jacobian to be able to compute approximate search directions reliably.

FORCES NLP provides different linear solvers to solve the symmetric indefinite form (11) and the normal equations form (13). In both cases a tailored structure-exploiting, block-wise Cholesky decomposition (Wright, 1993) with complexity $\mathcal{O}(N)$ can be used under the assumption that Φ and Y are positive definite. While in the quadratic convex case this assumption always holds for well-posed problems, the same does not hold for general nonlinear nonconvex problems. In the latter case, the Hessian of the Lagrangian is potentially indefinite and Y cannot be factorised with a Cholesky decomposition. Inertia correction techniques are used, for example in Wächter and Biegler (2006), to guarantee positive definiteness of the Hessian projected onto the null space of the constraint Jacobian. However, this approach requires several successive factorisations and regularisations of the KKT matrix (8), which lead to considerably increased computational demands. A common workaround used in nonconvex optimisation consists in approximating $\mathcal{H}(z, y_c)$ in order to guarantee its positive definiteness. In the following subsection several customised approaches for multistage programs implemented in the FORCES NLP package are presented and discussed.

3.2 Hessian approximation

3.2.1 Gauss-Newton

Consider a constrained nonlinear least-squares problem with cost function terms:

$$l_i(z_i) := \frac{1}{2} \|\eta_i(z_i)\|_2^2. \quad (15)$$

This is a rather common formulation which, in the MPC context, can encode, for example, regulation and tracking control objectives. The blocks of the Hessian of the

Lagrangian become

$$\begin{aligned} \mathcal{H}_i(z_{i+1}, z_i, y_{c,i}) &:= \sum_{k=1}^{r_{i-1}} y_{c,i-1[k]} \nabla_{z_i}^2 c_{i-1[k]}(z_i, z_{i-1}) \\ &+ \sum_{k=1}^{r_i} y_{c,i[k]} \nabla_{z_i}^2 c_{i[k]}(z_{i+1}, z_i) \\ &+ \nabla_{z_i} \eta_i(z_i) \nabla_{z_i} \eta_i(z_i)^T \\ &+ \sum_{k=1}^{v_i} (\eta_{i[k]}(z_i) \nabla_{z_i}^2 \eta_{i[k]}(z_i)), \end{aligned} \quad (16)$$

where v_i is the dimension of the residual η_i . Under the assumption of sufficiently small curvature of the functions η_i and c_i and small residuals $\eta_i(z_i)$, it is possible to neglect all the terms in (16) containing second-order derivatives, obtaining the following positive-semidefinite approximation:

$$\mathcal{H}_i(z_{i+1}, z_i, y_{c,i}) \approx \nabla \eta_i(z_i) \nabla \eta_i(z_i)^T \succeq 0 \quad (17)$$

and, assuming that the Jacobian of the residuals $\nabla \eta_i^T$ has full rank, a positive-definite approximation $\nabla \eta_i(z_i) \nabla \eta_i(z_i)^T \succ 0$ is obtained. This approach has been proven to work very well in practice in many cases (Diehl et al., 2009). The fact that the approximation is always positive-semidefinite allows one to exploit the efficient Cholesky decomposition proposed in Domahidi et al. (2012), while avoiding the evaluation of second-order derivatives for the Lagrangian, an operation that is potentially computationally demanding.

3.2.2 Block-wise Broyden-Fletcher-Goldfarb-Shanno Hessian updates

Whenever the problem does not take the form of a constrained nonlinear least-squares problem, or in case of highly nonlinear problems where the second derivatives cannot be neglected, it is necessary to use a different approximation.

A different approach is to use the high-rank block updates proposed in Bock and Plitt (1984). This second method consists of applying the low-rank Broyden-Fletcher-Goldfarb-Shanno (BFGS) update for each stage $i = 0, \dots, N$ separately. Given the updated primal and dual variables z_i^+ , $y_{c,i}^+$ and $y_{d,i}^+$, and the previous primal iterate z_i , the Hessian approximation B_i of the i th block (3.1) is updated as follows:

$$B_i \leftarrow B_i - \frac{B_i \delta \delta^T B_i}{\delta^T B_i \delta} + \frac{\gamma \gamma^T}{\delta^T \gamma}, \quad (18)$$

where

$$\begin{aligned}\delta &= z_i^+ - z_i \\ \gamma &= \nabla \mathcal{L}(z_i^+, y_{c,i}^+, y_{d,i}^+) - \nabla \mathcal{L}(z_i, y_{c,i}^+, y_{d,i}^+).\end{aligned}\quad (19)$$

The update (18) is guaranteed to preserve positive-definiteness of B_i if the condition $\delta^T \gamma > 0$ holds. In order to guarantee a positive-definite Hessian, it is possible to proceed in different ways. A first workaround is to skip the Hessian update whenever $\delta^T \gamma \leq 0$ and re-initialise the estimate whenever the update is skipped several times in succession. Such an approach is for example used in the IPOPT software package (Wächter & Biegler, 2006) and is also an option in FORCES NLP.

An alternative method is based on the modified update proposed by Powell (1978)

$$B_i \leftarrow B_i - \frac{B_i \delta \delta^T B_i}{\delta^T B_i \delta} + \frac{\omega \omega^T}{\delta^T \omega}, \quad (20)$$

with

$$\begin{aligned}\omega &:= \theta \gamma + (1 - \theta) B \delta \\ \theta &:= \begin{cases} 1 & \text{if } \delta^T \gamma \geq 0.2 \delta B \delta \\ \frac{0.8 \delta^T B \delta}{\delta^T B \delta - \delta^T \gamma} & \text{if } \delta^T \gamma < 0.2 \delta B \delta. \end{cases}\end{aligned}\quad (21)$$

It can be proved that the above update gives rise to super-linear convergence (Powell, 1978).

The initialisation of Hessian estimates B_i significantly affects the performance of the algorithm. For example, the common initialisation $B_i = \rho I$ with some a-priori chosen scalar $\rho > 0$ can be used. Alternatively, a user-defined initialisation can be provided.

3.3 Barrier strategy

In order to guarantee quick progress towards a local minimum, the KKT system (8) needs to be relaxed according to a given policy, commonly referred to as the *barrier strategy* which adjusts the barrier parameter τ in

$$\begin{aligned}r_S &= 0 \\ r_C &= 0 \\ r_D &= 0 \\ r_N &= \tau \mathbf{1}.\end{aligned}\quad (22)$$

While a Mehrotra predictor–corrector approach is used in Domahidi et al. (2012), due to its good empirical performance for convex quadratic problems, for nonlinear problems this heuristic can perform poorly in certain circumstances (Nocedal, Wächter, & Waltz, 2006).

A more robust approach is to monotonically decrease τ (Nocedal & Wright, 2006). Starting from a fixed initial value τ_0 , the relaxation is tightened whenever a solution that satisfies (22) to a given accuracy is found. In particular, given the current values of primal and dual variables and the barrier parameter, the optimality error is defined taking into account the maximum residuals in (8):

$$E_\tau := \max \{ \|r_S\|_\infty, \|r_C\|_\infty, \|r_D\|_\infty, \|r_N - \tau \mathbf{1}\|_\infty \} \quad (23)$$

and the barrier parameter is decreased whenever

$$E_\tau \leq \kappa_\epsilon \tau \quad (24)$$

for a given $\kappa_\epsilon > 0$. When condition (24) holds, the relaxation is updated according to

$$\tau \leftarrow \max \left\{ \frac{\epsilon_{\text{tol}}}{5}, \kappa_\tau \tau \right\}, \quad (25)$$

with constants $\kappa_\tau \in (0, 1)$ and ϵ_{tol} . When using such a barrier strategy, global convergence can be ensured, but slow convergence might be experienced (Nocedal et al., 2006). For this reason several accelerating adaptive strategies have been proposed in Nocedal et al. (2006), while global convergence is ensured by switching back to the monotone strategy when sufficient progress towards a solution cannot be made.

Different strategies have been analyzed and compared in Nocedal et al. (2006), based on variants of the Mehrotra predictor–corrector and other heuristics. In the following, an efficient approach (Vanderbei & Shanno, 1999) is reported that requires neglectable additional computational effort. This rule, used in the software package LOQO, determines the so-called *centering parameter* σ as

$$\sigma = 0.1 \min \left\{ 0.05 \frac{1 - \xi}{\xi}, 2 \right\}^3 \text{ with } \xi = \min_i \left\{ \frac{s_i y_{d,i}}{y_d^T s / q} \right\} \quad (26)$$

and the KKT relaxation is updated as follows:

$$\tau = \sigma \frac{s^T y_d}{q}. \quad (27)$$

The interested reader is referred to Nocedal et al. (2006) for more details on adaptive barrier strategies and their comparison. The FORCES NLP package implements several monotone and adaptive barrier strategies, including the ones described above.

3.4 Filter line-search method

In order to ensure convergence of the iterations, a backtracking line-search has been used based on the successful algorithm described in Wächter and Biegler (2006). Once the search direction has been computed, the maximum step lengths that maintain positivity in the slacks and multipliers are computed according to

$$\alpha_{\max} := \max\{\alpha \in (0, 1] : s + \alpha \Delta s \geq \tau s\}, \quad (28)$$

$$\alpha_l := \max\{\alpha \in (0, 1] : y_d + \alpha \Delta y_d \geq \tau y_d\}. \quad (29)$$

Then, a backtracking procedure is invoked, which iteratively reduces the primal step length from α_{\max} until a point is found that either improves feasibility or improves the barrier objective function:

$$\begin{aligned} \tilde{l}(\alpha_p) := & \sum_{i=0}^N l_i(z_i + \alpha_p \Delta z_i) \\ & - \tau \sum_{i=0}^N \sum_{k=1}^{r_i} \log(s_{i[k]} + \alpha_p \Delta s_{i[k]}). \end{aligned}$$

For more details, see Wächter and Biegler (2006).

3.5 Integration and sensitivity generation

Building the linear system (8) associated with the computation of the Newton steps requires integration of differential equations and computation of the sensitivities associated with the obtained trajectories. In order to do so, explicit or implicit integration methods can be chosen depending on several parameters such as stiffness of the differential equation, computational burden associated with function evaluations and number of states.

Integration and sensitivity generation routines can be computationally expensive and their careful implementation is of crucial importance in order to reduce computation times. While explicit schemes might be appropriate for non-stiff systems and when function evaluations can be carried out cheaply, in the presence of stiff dynamics or expensive function evaluations it is generally necessary to use implicit schemes. The latter require the solution of implicit nonlinear equations that can be rather expensive.

FORCES NLP implements both explicit integrators and the state-of-the-art algorithms for implicit Runge-Kutta schemes introduced in Quirynen, Vukov, Zanon, and Diehl (2014), that have been proven to be orders of magnitude faster than commercial packages for problems

of small to medium size. Such schemes exploit Newton-type iterations with a fixed Jacobian of the Runge-Kutta equations. In this way, a single factorisation is carried out per integration step for both integration and sensitivity generation resulting in a highly efficient algorithm. A detailed description and benchmarking of the scheme can be found in Quirynen et al. (2014).

3.6 Initialisation and warm-start

In order to initialise the algorithm, the user-provided initial point z^0 is modified such that it is feasible, i.e. ensuring that bounds on the primal variables are satisfied and that positivity constraints on slacks and bounds multipliers are satisfied. The multipliers $y_{c,i}$ associated with bounds are initialised to one componentwise and a least-squares estimate for the equality multipliers $y_{d,i}$ is computed using the possibly modified initial point (Wächter & Biegler, 2006).

Notice that, in the context of nonlinear MPC, where a series of neighbouring problems have to be solved as the initial conditions change, SQP methods can exploit warm-starts available from the solution of previous instances. This feature can lead to a significant reduction in the number of iterations required. In particular, for the RTI scheme theoretical results are present in the literature that guarantee local convergence if an initial guess that is “sufficiently” close to the optimal solution is available Diehl et al. (2007).

Warm-start of interior-point methods is less straightforward. Issues and potential solutions associated with the warm-start of interior-point methods are discussed in the literature for linear (Yildirim & Wright, 2002), quadratic (Gondzio & Grothey, 2006) and nonlinear nonconvex programs (Benson & Shanno, 2008; Forsgren, 2006). Moreover, algorithms tailored to MPC applications have been proposed in Shahzad and Goulart (2011), Shahzad, Kerrigan, and Constantinides (2010), and Zanelli, Quirynen, Jerez, and Diehl (2017). The current implementation of FORCES NLP does not exploit warm-starts, but the warm-starting techniques available in the literature can be in principle applied to the algorithm implemented.

4. Numerical results

The algorithm described has been implemented in the software package for embedded optimisation FORCES NLP as an extension of the previously existing primal-dual interior-point solver for convex QCQPs (Domahidi et al., 2012). The tool consists of a code generation engine

capable of providing efficient, statically allocated, standalone C code with a small memory footprint.

In this section, the performance achieved with FORCES NLP (Domahidi & Jerez, 2016) is compared with IPOPT (Wächter & Biegler, 2006) and the solvers generated with the ACADO Toolkit (Houska et al., 2011). ACADO is a software environment for automatic control and dynamic optimisation that provides code generation functionalities to export C code that implements the RTI algorithm (Diehl, Bock, & Schlöder, 2005). While IPOPT can, in general, more reliably solve NLPs to local optimality, it often leads to long computation times, the RTI-based solver in ACADO has been designed to provide approximate solutions within shorter computation times. By using FORCES NLP it is possible, for the examples presented in the following, to combine the same closed-loop performance as the one obtained by IPOPT, with computation times in the range of the ones achievable with the RTI scheme, hence offering a better trade-off between control performance and computational delay.

In order to compare the three solvers, tracking nonlinear optimal control problems of the form

$$\begin{aligned} \min_{\substack{x_0, \dots, x_N \\ u_0, \dots, u_{N-1}}} & \sum_{i=0}^{N-1} l_i(x_i, u_i) + l_N(x_N) \\ \text{s.t. } & x_{i+1} = c(x_i, u_i) \quad \forall i = 0, \dots, N-1 \\ & x_{\min} \leq x_i \leq x_{\max} \quad \forall i = 0, \dots, N \\ & u_{\min} \leq u_i \leq u_{\max} \quad \forall i = 0, \dots, N-1, \end{aligned}$$

will be formulated and solved for three different nonlinear systems. The function c describes the discretised dynamics obtained using an explicit Runge-Kutta integrator of order four (RK4). The simulations in the following have been done using Gauss-Newton RTI solvers code generated with ACADO using two different QP solvers: the active-set solver qpOASES (Ferreau et al., 2014) and the interior-point solver FORCES PRO (Domahidi et al., 2012).

FORCES NLP has been configured to use block-wise BFGS Hessian approximations and a LOQO adaptive barrier strategy (Nocedal et al., 2006). Routines that evaluate functions and their derivatives required for integration, sensitivity generation and the BFGS updates have been code generated using CasADi (Andersson, Akesson, & Diehl, 2012). For IPOPT, the CasADi interface has been used to set up the optimal control formulations. As compiling the code generated by CasADi for functions, Jacobians and Hessians evaluations resulted in *out-of-memory* errors for some of the examples, the routines have been




Figure 1. Robotic manipulator with two links and two joints (Siciliano et al., 2009).

compiled and timed separately in order to have a fair estimate of the total execution time for IPOPT. Results with both exact and BFGS approximate Hessians are reported and the sparse direct linear solver MA57 (Duff, 2004) is used. For FORCES NLP and IPOPT, the default termination conditions have been used. All the benchmarks have been run on a laptop XPS13 9350 with an Intel i7-6500U CPU (maximum boost frequency of 3.10 GHz) and Ubuntu 15.10 as operating system. All the examples described in the next sections are available online at <https://github.com/embotech/forcesnlp-examples>.

4.1 Robotic manipulator

Consider the robotic manipulator with two links and two revolute joints (Siciliano, Sciavicco, Villani, & Oriolo, 2009) illustrated in Figure 1. Two electric drives apply torques to the joints in order to steer the end-effector to the desired position.

The dynamical model of such a system can be derived with the Lagrangian approach:

$$\begin{aligned} \ddot{\theta}_1 &= \gamma_1 \\ \ddot{\theta}_2 &= \gamma_2 \\ \dot{\tau}_1 &= \tau_{1r} \\ \dot{\tau}_2 &= \tau_{2r}, \end{aligned} \tag{30}$$

with

$$\begin{aligned} \gamma_1 &= \frac{1}{(\alpha_1 - \alpha_2 \frac{\beta_1}{\beta_2}) \beta_2} (\beta_4 + \beta_3 \dot{\theta}_1^2 - \tau_2) - \alpha_3 \dot{\theta}_1 \dot{\theta}_2 + \\ &\quad - \alpha_4 \dot{\theta}_2 - \alpha_5 + \tau_1 \\ \gamma_2 &= \frac{1}{\beta_2} (\tau_2 - \beta_1 \gamma - \beta_3 \dot{\theta}_1^2 - \beta_4), \end{aligned}$$

Table 1. Robotic manipulator: parameter values.

Parameter	Value	Description
$m_{M_1} = m_{M_2}$	0.3 kg	Motor mass
$k_{r_1} = k_{r_2}$	1.0	Gear reduction ratio
$r_{M_1} = r_{M_2}$	0.05 m	Motor radius
$I_{M_1} = I_{M_2}$	$\frac{1}{2} m_{M_1} r_{M_1}^2$	Motor moment of inertia
$m_{l_1} = m_{l_2}$	3.0 kg	Link mass
$I_{l_1} = I_{l_2}$	$\frac{1}{12} m_{l_1} l_1^2$	Link moment of inertia
$a_1 = a_2$	$\frac{1}{2} l_1$	Link centre of mass relative position
$l_1 = l_2$	1 m	Link length
g	9.81 m/s ²	Gravitational acceleration

and

$$\begin{aligned}
\alpha_1 &= I_{l_1} + m_{l_1} l_1^2 + k_{r_1}^2 I_{m_1} + I_{l_2} + m_{l_2} (a_1^2 + l_2^2 \\
&\quad + 2a_1 l_2 \cos(\theta_2)) + I_{m_2} + m_{M_2} a_1^2) \\
\alpha_2 &= I_{l_2} + m_{l_2} (l_2^2 + a_1 l_2 \cos(\theta_2)) + k_{r_2} I_{m_2} \\
\alpha_3 &= -2.0 m_{l_2} a_1 l_2 \sin(\theta_2) \\
\alpha_4 &= -m_{l_2} a_1 l_2 \sin(\theta_2) \\
\alpha_5 &= (m_{l_1} l_1 + m_{M_2} a_1 + m_{l_2} a_1) g \cos(\theta_1) \\
&\quad + m_{l_2} l_2 g \cos(\theta_1 + \theta_2) \\
\beta_1 &= I_{l_2} + m_{l_2} (l_2^2 + a_1 l_2 \cos(\theta_2)) + k_{r_2} I_{m_2} \\
\beta_2 &= I_{l_2} + m_{l_2}^2 + k_{r_2} k_{r_2} I_{m_2} \\
\beta_3 &= m_{l_2} a_1 l_2 \sin(\theta_2) \\
\beta_4 &= m_{l_2} l_2 g \cos(\theta_1 + \theta_2),
\end{aligned}$$

where θ_1, θ_2 are the angles describing the configurations of the manipulator and τ_1, τ_2 are the torques applied to the joints. The numerical values of the fixed geometrical and mechanical parameters used for the simulations are reported in Table 1.

In order to formulate the optimal control problem, states $x := [\theta_1, \dot{\theta}_1, \theta_2, \dot{\theta}_2, \tau_1, \tau_2]^T$ and controls $u := [\tau_{1r}, \tau_{2r}]^T$ are defined and the following constraints are imposed:

$$\begin{aligned}
-100 \text{ Nm} &\leq \tau_1 \leq 70 \text{ Nm} \\
-100 \text{ Nm} &\leq \tau_2 \leq 70 \text{ Nm} \\
-200 \text{ Nm/s} &\leq \tau_{1r} \leq 200 \text{ Nm/s} \\
-200 \text{ Nm/s} &\leq \tau_{2r} \leq 200 \text{ Nm/s}.
\end{aligned}$$

A quadratic cost function is used that penalises the deviation of states and inputs from a reference:

$$\begin{aligned}
l_i &:= (x - x_r)^T Q (x - x_r) + u^T R u \\
l_N &:= (x_N - x_r)^T Q_N (x_N - x_r),
\end{aligned}$$

with

$$\begin{aligned}
Q &= Q_N := \text{diag}(10^3, 10^{-1}, 10^3, 10^{-1}, 10^{-2}, 10^{-2}) \\
R &:= \text{diag}(10^{-2}, 10^{-2})
\end{aligned}$$

and reference defined as

$$x_r = \begin{cases} t < 10s & [1.2, 0, -1.2, 0, 0, 0] \\ t \geq 10s & -[1.2, 0, -1.2, 0, 0, 0]. \end{cases} \quad (31)$$

A prediction horizon of $T_f = 2$ s over $N = 20$ shooting nodes is used with two intermediate RK4 integration steps per shooting interval.

Figure 2 shows the closed-loop trajectories obtained. In this case the difference between performance achieved by ACADO and FORCES NLP is minor and the relative closed-loop suboptimality is 0.85%. The same local minima are obtained with both FORCES NLP and IPOPT at each sampling instant. The same holds for the two QP solvers used, FORCES PRO and qpOASES. For this reason, only the trajectories obtained with qpOASES are shown in the plots.

Figure 3 compares the number of iterations and computation times for the three solvers. ACADO and FORCES NLP achieve the fastest computation times, with FORCES NLP being only about two to three times slower than ACADO with FORCES PRO as QP solver. In FORCES NLP 0.0286 ms per iteration are spent in the integrators on average, in order to compute state trajectories and sensitivities.

It was not possible to compile the generated C code for the evaluations of functions, Jacobians and Hessians for CasADi-IPOPT, due to *out-of-memory* errors. For this reason, these routines have been code generated, compiled and timed separately and the total execution time has been estimated.

4.2 Quadcopter

In this section, a tracking formulation is considered with the goal of stabilising a quadcopter around its hover state. The control of such a system with nonlinear MPC has been previously addressed in Quirynen et al. (2013a), where the RTI algorithm is used to obtain approximate solutions. In the following, it will be shown how the performance can be significantly improved when solving the problems to a local minimum at every sampling instant with moderate additional computational burden using the solver introduced. Consider the following model of




Figure 2. Robotic manipulator tracking a time-varying reference. The figures compare the closed-loop trajectories obtained with the RTI scheme with shifts (`'shifting.strategy=1'`) (dashed) and FORCES NLP (solid). ACADO uses qpOASES (with warm-starting) as QP solver. In this case, there is a minor deviation in the closed-loop trajectories and the relative suboptimality is 0.85%. The trajectories obtained with IPOPT and ACADO with FORCES PRO as QP solver are not shown in the plot as they are identical to the ones obtained with FORCES NLP and ACADO with qpOASES, respectively.

a quadcopter adapted from Quirynen et al. (2013a):

$$\begin{aligned}
 \dot{p} &= v \\
 \dot{q} &= \frac{1}{2} E^T \Omega \\
 \dot{v} &= \frac{1}{m} R F - g \mathbf{1}_z \\
 \dot{\Omega} &= J^{-1} (T + \Omega \times J \Omega) \\
 \dot{\omega} &= \omega_r,
 \end{aligned} \tag{32}$$

where p and v represent position and velocity of the quadcopter respectively and q and Ω are its orientation expressed in quaternion representation and its angular velocity. It is assumed that angular accelerations of the propellers ω_r can be tracked instantaneously, hence they are considered as inputs to the system. The constants J and m are the inertia matrix of the vehicle and mass, respectively, while the forces and torques applied to the

system are described by

$$\begin{aligned}
 F &= \sum_{k=1}^4 \frac{1}{2} \rho A C_L \omega_k^2 \mathbf{1}_z \\
 T &= [T_1 \ T_2 \ T_3]^T,
 \end{aligned} \tag{33}$$

with

$$\begin{aligned}
 T_1 &:= \frac{A C_l L \rho (\omega_2^2 - \omega_4^2)}{2}, & T_2 &:= \frac{A C_l L \rho (\omega_1^2 - \omega_3^2)}{2}, \\
 T_3 &:= \frac{A C_d L \rho (\omega_1^2 - \omega_2^2 + \omega_3^2 - \omega_4^2)}{2},
 \end{aligned}$$

where ρ is the air density, C_D and C_L are the drag and lift coefficients and A is the area of the propellers. The values of the parameters appearing in the model are listed in Table 2.

States and controls are defined as $x := [p^T, q^T, v^T, \Omega^T, \omega^T]^T$ and $u := \omega_r$ and the following

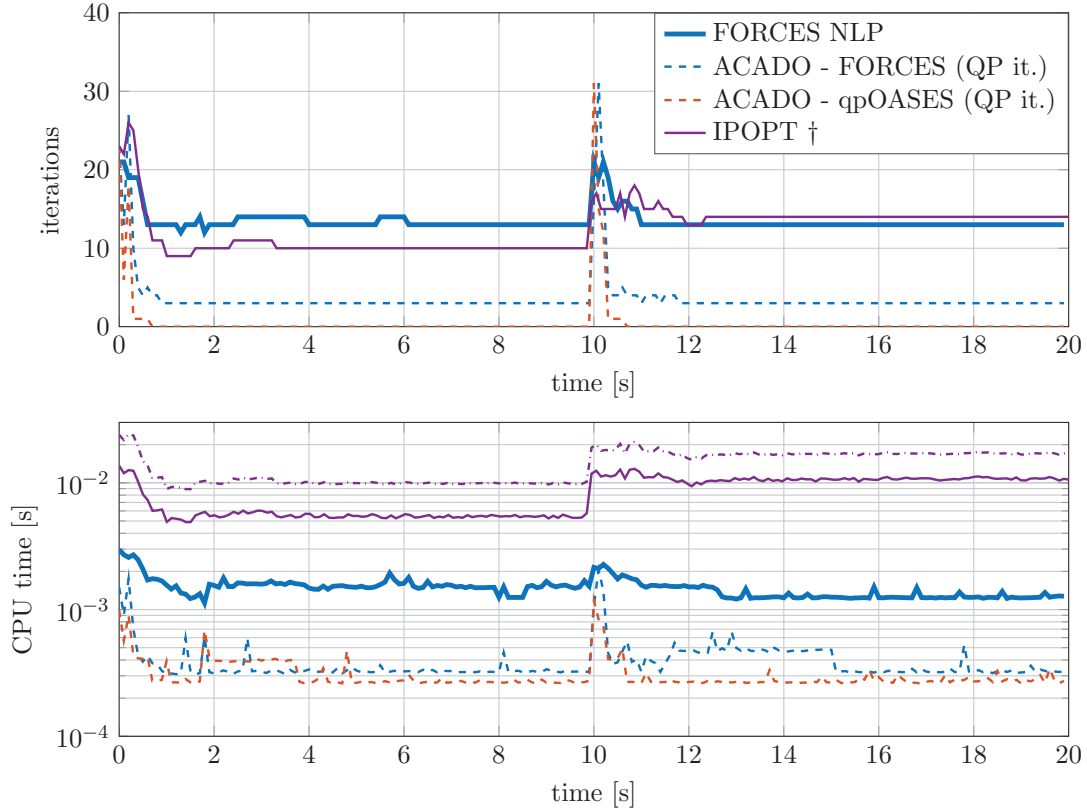


Figure 3. Number of iterations and computation times for all solvers. Interior-point iterations are reported for IPOPT and FORCES NLP, while the number of QP iterations to solve every QP are reported for ACADO. ACADO exploits shifting ('shifting.strategy=1') and qpOASES is warm-started. ACADO is always faster than FORCES NLP and it achieves a worst-case execution time that is less than two to three times shorter. While both nonlinear interior-point solvers solve the OCPs to a local minimum, FORCES NLP has a worst-case execution time that is about eight times shorter than the one obtained with IPOPT. Integration and sensitivity generation in FORCES NLP requires, for this example, 0.0286 ms per iteration on average. † - It was not possible to compile the code generated by CasADi-IPOPT for functions, Jacobians and Hessians evaluations, due to out-of-memory errors. These routines have been code generated, compiled and timed separately in order to estimate CPU time spent performing these operations (total time estimated in dash-dotted style).

constraints are imposed:

$$\begin{aligned}
 20 \text{ rad/s} &\leq \omega_k \leq 50 \text{ rad/s}, & \forall k = 1, \dots, 4 \\
 -40 \text{ rad/s}^2 &\leq \omega_{rk} \leq 40 \text{ rad/s}^2, & \forall k = 1, \dots, 4 \\
 -40 \text{ rad/s} &\leq \Omega_k \leq 40 \text{ rad/s}, & \forall k = 1, \dots, 3 \\
 -10 \text{ m} &\leq p_k \leq 10 \text{ m}, & \forall k = 1, \dots, 3 \\
 -100 \text{ m/s} &\leq v_k \leq 100 \text{ m/s}, & \forall k = 1, \dots, 3.
 \end{aligned}$$

A prediction horizon of $T_f = 1$ s over $N = 30$ shooting nodes is used with a single RK4 integration step per

shooting interval. A strictly convex quadratic cost function and terminal cost have been used to penalise deviation of states and inputs from a given reference:

$$\begin{aligned}
 l_i &:= (x_i - x_r)^T Q (x_i - x_r) + (u_i - u_r)^T R (u_i - u_r) \\
 l_N &:= (x_N - x_{rN})^T Q_T (x_N - x_{rN}), \\
 Q_N = Q &= \text{diag}(\mathbf{1}_3^T \cdot 50, \mathbf{1}_4^T \cdot 20, \mathbf{1}_3^T \cdot 5, \mathbf{1}_3^T \cdot 5, \\
 &\quad \mathbf{1}_4^T \cdot 0.01) \\
 R &= \text{diag}(\mathbf{1}_3 \cdot 5)
 \end{aligned} \tag{34}$$

with constant references $x_r := [\mathbf{0}_3^T, 1, 0, 0, 0, \mathbf{0}_3^T, \mathbf{0}_3^T, \bar{\omega} \mathbf{1}_4^T]$ and $u_r := [0, 0, 0, 0]$, where $\bar{\omega} = 39.939$ rad/s.

Figure 4 shows the resulting closed-loop trajectories when using $N = 30$ shooting nodes. Both controllers manage to stabilise the system, steering the states to the specified reference. However, by computing a local minimum at every feedback step, the nonlinear interior-point solver achieves better control performance. The

Table 2. Quadcopter: parameter values.

Parameter	Value	Description
ρ	1.225 kg/m ³	Air density
A	0.1 m ²	Propeller area
C_l	0.25	Lift coefficient
C_d	$0.3 \cdot C_l$	Drag coefficient
m	10 kg	Quadcopter mass
g	9.81 m/s ²	Gravitational acceleration
$J_1 = J_2 = J_3$	0.25 kg · m/s ²	Link length

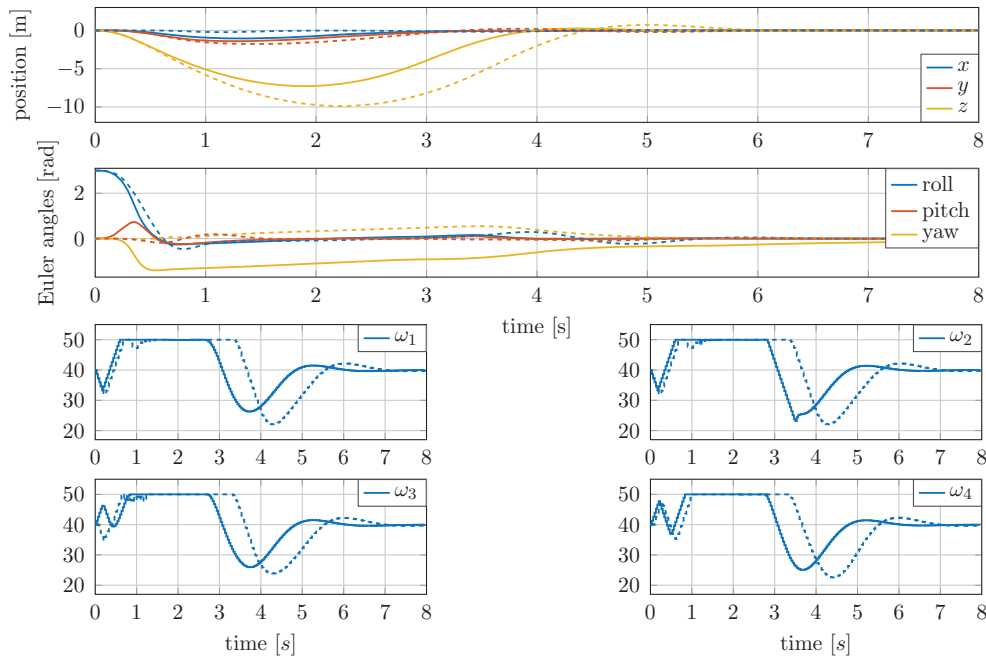


Figure 4. Stabilisation of the quadcopter around hovering state: closed-loop trajectories obtained with FORCES NLP (solid) and ACADO with qpOASES (dashed) for $N = 30$. Although different trajectories are obtained with IPOPT, the latter are not reported for clarity of the plots. ACADO uses warm-starts and shifts to initialise subsequent subproblems. The nonlinear interior-point solver computes a local minimum at every feedback step resulting in an improved performance. Relative closed-loop suboptimality associated with the approximate scheme is 69.07%.

closed-loop trajectories obtained with IPOPT differ from the one obtained with FORCES NLP, due to the fact that the solver computes different local minima. However, the trajectories computed by IPOPT are not reported for clarity of the plots and are not taken into account in the suboptimality comparison in Table 3 as both solvers find locally optimal trajectories. The closed-loop cost achieved with FORCES NLP is taken as a reference to compute the suboptimality of the RTI scheme.

In order to improve the closed-loop performance obtained with the RTI, it is possible to perform multiple SQP steps per OCP instance. Every iteration would require an additional linearisation and QP solution. Table 3 shows the obtained computation times and suboptimality for the OCP with an increasing number of shooting nodes. When applying several iterations to the same instance, shifts are applied only when updating the initial condition.

Table 3. Quadcopter example: worst-case CPU time and closed-loop relative suboptimality for increasing horizon lengths. For ACADO, the minimum number of SQP iterations required to stabilise the system is taken into account (in brackets next to timings). For small problems, ACADO – qpOASES can achieve moderate suboptimality with considerably shorter computation times. Although FORCES NLP is always slower than ACADO, when increasing the horizon length, rather competitive computation times can be achieved with improved performance with respect to the RTI algorithm.

	$N = 10$		$N = 20$		$N = 30$	
	Time [ms] (it.)	Sub. %	time [ms] (it.)	Sub. %	Time [ms] (it.)	Sub. %
RTI – qpOASES	5.09 (3)	11.86	24.43 (2)	41.19	42.72(1)	69.07
RTI – FORCES ^a	15.52(3)	20.27	35.76(3)	11.13	44.57(2)	10.47
FORCES NLP	18.17	–	45.18	–	51.71	–
IPOPT ^b	109.7	–	283.1	–	494.6	–

^aNotice that, when using FORCES as QP solver, it is necessary to damp the Newton steps in order to avoid ill-conditioned QPs which would make the solver fail. This justifies the different suboptimality incurred with respect to ACADO – qpOASES.

^bComputation times for IPOPT are estimated based on separate timings of code generated function evaluations.

Table 4. Solve times and number of iterations for the vehicle path planning problem. Timings with and without function evaluations are reported. The routines evaluating functions and their first- and second-order derivatives have been code generated using the CasADi interface to IPOPT. Notice that, especially when computing exact Hessians, the time spent in function evaluations can be significant.

	Time w/o f. evals	Total solution time	Iterations
FORCES NLP - BFGS	7.0 ms	8.0 ms	74
IPOPT - BFGS	95 ms	140 ms	70
IPOPT - exact Hessian	105 ms	258 ms	117

ACADO is always faster than FORCES NLP, but, especially for problems with longer horizons, the worst-case execution time is only moderately increased. A possible reason for this is that, for this example, the time spent in solving the ODE and computing sensitivities is a rather small fraction of the time required to carry out an interior-point iteration.

As it was not possible to compile the code generated by the interface to IPOPT in CasADi for function evaluations, the computation times reported in Table 3 are estimates computed based on separate timings for function evaluations. Routines for the evaluation of functions, Jacobians and Hessians have been code generated in CasADi, compiled and timed separately. When using FORCES NLP, a considerable speedup can be achieved with respect to timings estimated for IPOPT.

4.3 Vehicle path planning

Real-time trajectory generation is an ubiquitous task for next generation robotics and autonomous driving applications. An optimisation-based approach provides a systematic way for computing optimal trajectories with respect to the performance goals of the application. However, the resulting optimisation problems are often challenging due to nonlinear models and nonconvex constraints, for instance, when avoiding obstacles.

Consider the following simple nonlinear vehicle model:

$$\begin{aligned}\dot{y} &= v \cos(\theta) \\ \dot{z} &= v \sin(\theta) \\ \dot{v} &= F/m \\ \dot{\theta} &= s/I,\end{aligned}$$

where y and z are the Cartesian coordinates of the vehicle, v is the linear velocity and θ is the heading angle. There are two control inputs to the model: the acceleration force F and the steering torque s . The vehicle mass is $m = 1\text{ kg}$ and its moment of inertia is $I = 1\text{ kgm}^2$.

For this example we formulate a trajectory generation problem where the objective is to maximise progress on the z -direction, while minimising the applied force and torque, i.e.

$$l_i := -100z + 0.1F^2 + 0.001s^2.$$

The dynamics (35) are discretised with the explicit RK4 integrator using $N = 50$ shooting nodes over a prediction horizon of 5 s. The manoeuvre is subject to a set of constraints, involving both the simple bounds

$$\begin{aligned}-5\text{ N} &\leq F \leq 5\text{ N} \\ -1\text{ Nm} &\leq s \leq 1\text{ Nm} \\ -3\text{ m} &\leq y \leq 0\text{ m} \\ 0\text{ m} &\leq z \leq 3\text{ m} \\ 0\text{ m/s} &\leq v \leq 2\text{ m/s} \\ 0\text{ rad} &\leq \theta \leq \pi\text{ rad},\end{aligned}$$

as well as the nonlinear nonconvex constraints representing obstacles

$$\begin{aligned}1\text{ m}^2 &\leq y^2 + z^2 \leq 9\text{ m}^2 \\ 0.9025\text{ m}^2 &\leq (y + 2)^2 + (z - 2.5)^2.\end{aligned}$$

In addition, the vehicle should be at standstill and with a heading angle of zero degrees at the end of the manoeuvre. The simulation using IPOPT exploits code generated C routines for the evaluations of functions and their first- and second-order derivatives. The generated code has been compiled with optimisation flag `-O3`. Figure 5 shows the computed optimal trajectories, which are the same for FORCES NLP and IPOPT. However, there is a large difference in the time required to compute the solution with the two solvers. Table 4 shows that FORCES NLP is more than one order of magnitude faster than IPOPT for this trajectory generation problem.

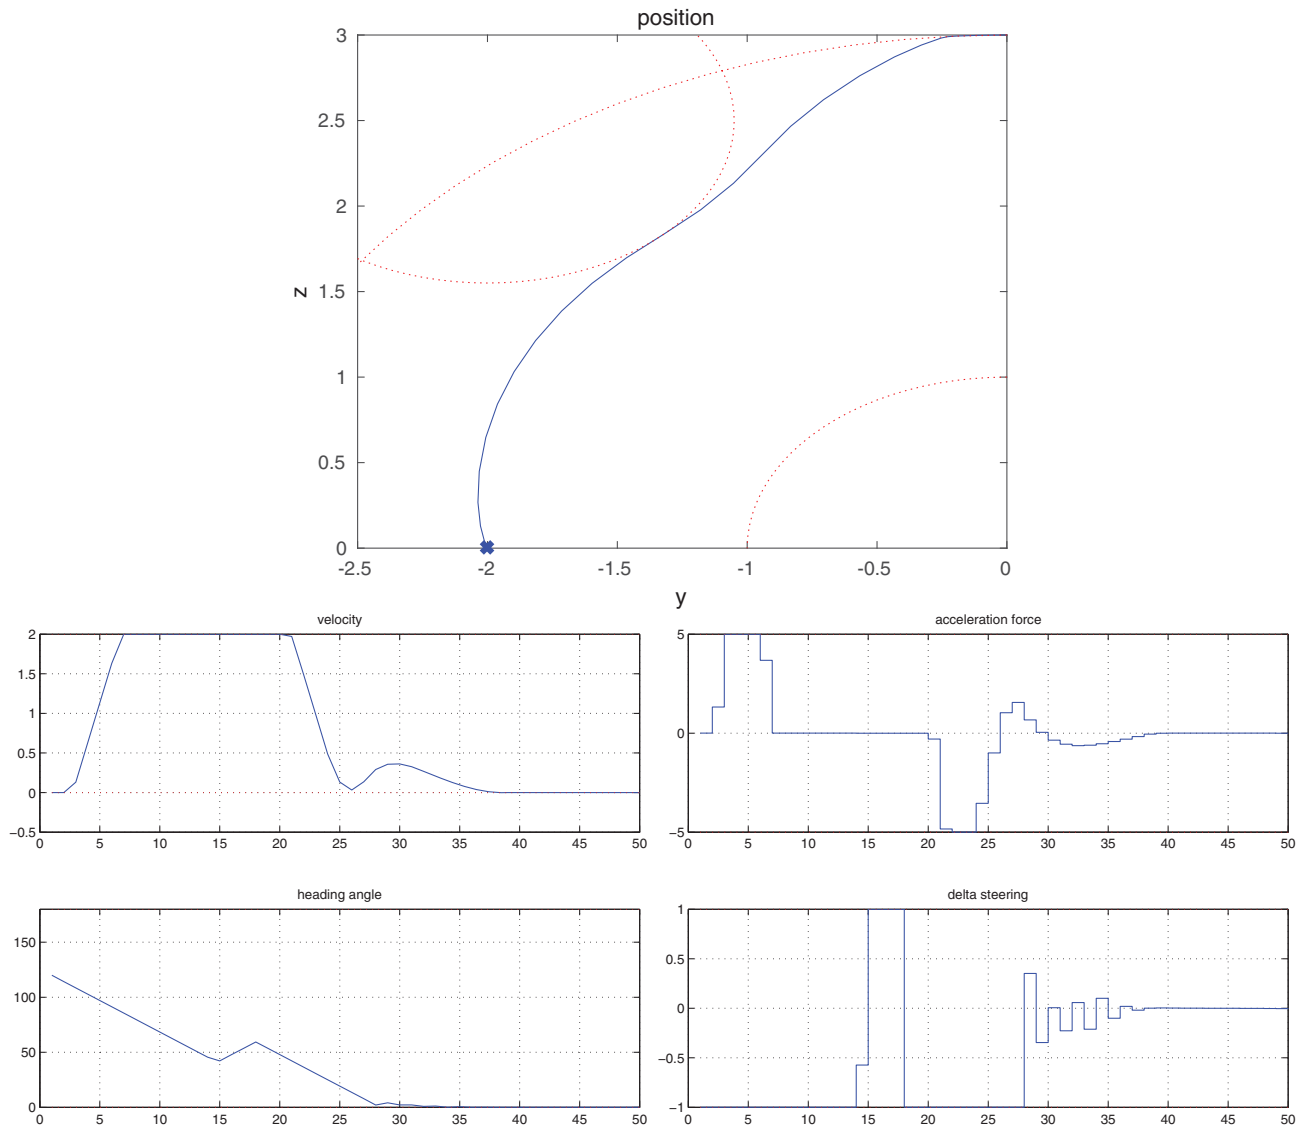


Figure 5. Trajectory of the vehicle on a two-dimensional space (top), vehicle's velocity and heading angle (left), and actuators (right). Constraints are represented with dotted red lines. The vehicle starts at position (-2,0) and should maximise progress in z-direction, while being at standstill and at heading angle of zero degrees at its final position after 50 integration steps.

5. Conclusions

The software package FORCES NLP that provides an efficient tailored implementation of an NLP solver for nonlinear nonconvex multistage problems has been presented. The package implements a structure-exploiting nonlinear interior-point method with approximate Hessians and generates standalone efficient C code that can be readily deployed on any embedded platform. The performance of the solver has been assessed on two optimal control and trajectory generation examples. It has been shown that, for examples for which it makes a difference to solve NLPs to local optimality, computation times in the same range as with fast approximate schemes can be achieved, while considerable improvements in the control performance can be attained.

Moreover, the presented software has been compared to the state-of-the-art code for nonlinear optimisation IPOPT, achieving speedups of up to an order of magnitude. The numerical examples comparing FORCES NLP, ACADO and IPOPT used in this paper are all available on <https://github.com/embotech/forcesnlp-examples>.

Disclosure statement

No potential conflict of interest was reported by the authors.

References

- Andersson, J., Akesson, J., & Diehl, M. (2012). CasADi – A symbolic package for automatic differentiation and optimal control. In S. Forth, P. Hovland, E. Phipps, J. Utke, & A. Walther (Eds.), *Recent Advances in Algorithmic*

- Differentiation, Lecture Notes in Computational Science and Engineering* (Vol. 87, pp. 297–307). Berlin: Springer.
- Bemporad, A., & Morari, M. (1999). Control of systems integrating logic, dynamics, and constraints. *Automatica*, 35(3), 407–427.
- Benson, H., & Shanno, D. (2008). Interior-point methods for nonconvex nonlinear programming: Regularization and warmstarts. *Computational Optimization and Applications*, 40(143), 143–189.
- Bock, H., & Plitt, K. (1984). A multiple shooting algorithm for direct solution of optimal control problems. In *Proceedings of the 9th IFAC World Congress* (pp. 242–247). Budapest: Pergamon Press.
- Chen, W., Balance, D., & O'Reilly, J. (2000). Model predictive control of nonlinear systems: Computational burden and stability. *IEEE Proceedings – Control Theory and Applications*, 147(4), 387–394.
- Deng, J., Becerra, V. M., & Stobart, R. (2009). Input constrained handling in an MPC/feedback linearization scheme. *International Journal of Applied Mathematics and Computer Science*, 19(2), 219–232.
- Diehl, M., Bock, H. G., & Schlöder, J. P. (2005). A real-time iteration scheme for nonlinear optimization in optimal feedback control. *SIAM Journal on Control and Optimization*, 43(5), 1714–1736.
- Diehl, M., Bock, H. G., Schlöder, J. P., Findeisen, R., Nagy, Z., & Allgöwer, F. (2002). Real-time optimization and nonlinear model predictive control of processes governed by differential-algebraic equations. *Journal of Process Control*, 12(4), 577–585.
- Diehl, M., Ferreau, H., & Haverbeke, N. (2009). Nonlinear model predictive control, efficient numerical methods for nonlinear MPC and moving horizon estimation. *Lecture Notes in Control and Information Sciences*, 384, 391–417.
- Diehl, M., Findeisen, R., & Allgöwer, F. (2007). A stabilizing real-time implementation of nonlinear model predictive control. In L. Biegler, O. Ghattas, M. Heinkenschloss, D. Keyes & B. van Bloemen Waanders (Eds.), *Real-Time and Online PDE-Constrained Optimization*, 23–52. SIAM.
- Domahidi, A., & Jerez, J. (2016). *FORCES Nonlinear Programming Solver*. embotech GmbH. Retrieved from <https://www.embotech.com/FORCES-Pro/User-Manual/High-level-Interface>. Zurich, Switzerland.
- Domahidi, A., Zraggen, A., Zeilinger, M., Morari, M., & Jones, C. (2012). Efficient interior point methods for multistage problems arising in receding horizon control. In *IEEE Conference on Decision and Control* (pp. 668–674). Maui, HI.
- Duff, I. (2004). Ma57 – a code for the solution of sparse symmetric definite and indefinite systems. *ACM Transactions on Mathematical Software*, 30(2), 118–144.
- Falcone, P., Borrelli, F., Asgari, J., Tseng, H. E., & Hrovat, D. (2007). Predictive active steering control for autonomous vehicle systems. *IEEE Transactions on Control Systems Technology*, 15(3), 566–580.
- Ferreau, H., Houska, B., Geibelen, K., & Diehl, M. (2011). Real-time control of a kite-model using an auto-generated nonlinear MPC algorithm. *Proceedings of 18th IFAC World Congress*, 44(1), 2488–2493.
- Ferreau, H., Kirches, C., Potschka, A., Bock, H., & Diehl, M. (2014). qpOASES: A parametric active-set algorithm for quadratic programming. *Mathematical Programming Computation*, 6(4), 327–363.
- Findeisen, R., & Allgöwer, F. (2004). Computational delay in nonlinear model predictive control. In *Proceedings of International Symposium on Advanced Control of Chemical Processes* (pp. 427–432). Hong Kong.
- Forsgren, A. (2006). *On warm starts for interior methods* (pp. 51–66). Boston, MA: Springer.
- Frison, G., Sorensen, H. B., Dammann, B., & Jørgensen, J. B. (2014). High-performance small-scale solvers for linear model predictive control. In *Proceedings of European Control Conference* (pp. 128–133). Strasbourg, France.
- Garcia, C., Prett, D., & Morari, M. (1988). Model predictive control: Theory and practice – A survey. *Automatica*, 25(3), 335–348.
- Giselsson, P. (2012). Execution time certification for gradient-based optimization in model predictive control. In *Proceedings of the 51st IEEE Conference on Decision and Control* (pp. 3165–3170). Maui, HI.
- Gondzio, J., & Grothey, A. (2006). A new unblocking technique to warmstart interior point methods based on sensitivity analysis. *SIAM Journal on Control and Optimization*, 19(3), 1184–1210.
- Houska, B., Ferreau, H., & Diehl, M. (2011). An auto-generated real-time iteration algorithm for nonlinear MPC in the microsecond range. *Automatica*, 47(10), 2279–2285.
- Jerez, J., Goulart, P., Richter, S., Constantinides, G., Kerrigan, E., & Morari, M. (2014). Embedded online optimization for model predictive control at megahertz rates. *IEEE Transactions on Automatic Control*, 59(12), 3238–3251.
- Kang, J., Cao, Y., Word, D., & Laird, C. (2014). An interior-point method for efficient solution of block-structured NLP problems using an implicit Schur-complement decomposition. *Computers & Chemical Engineering*, 71, 563–573.
- Kothare, M. V., Balakrishnan, V., & Morari, M. (1996). Robust constrained model predictive control using linear matrix inequalities. *Automatica*, 32(10), 1361–1379.
- Nicolao, G. D., Magni, L., & Scattolini, R. (1998). Nonlinear model predictive control, Stability and robustness of nonlinear receding horizon control. *Progress in Systems and Control Theory*, 3–22.
- Nocedal, J., Wächter, A., & Waltz, R. (2006). Adaptive barrier update strategies for nonlinear interior methods. *SIAM Journal on Optimization*, 19(4), 1674–1693.
- Nocedal, J., & Wright, S. J. (2006). *Numerical optimization. Springer Series in Operations Research and Financial Engineering* (2nd ed.). Berlin: Springer.
- Ohtsuka, T. (2004). A continuation/GMRES method for fast computation of nonlinear receding horizon control. *Automatica*, 40(4), 563–574.
- Powell, M. (1978). A fast algorithm for nonlinearly constrained optimization calculations. In G. A. Watson (Ed.), *Numerical Analysis* (pp. 144–157). Berlin: Springer.
- Qin, S. J., & Badgwell, T. A. (2003). A survey of industrial model predictive control technology. *Control Engineering Practice*, 11(7), 733–764.
- Quirynen, R., Gros, S., & Diehl, M. (2013a). Efficient NMPC for nonlinear models with linear subsystems. In *Proceedings of the 12th IEEE Conference on Decision and Control* (pp. 5101–5106). Florence, Italy.
- Quirynen, R., Gros, S., & Diehl, M. (2013b). Fast auto generated ACADO integrators and application to MHE

- with multi-rate measurements. In *Proceedings of the 12th European Control Conference* (pp. 3077–3082). Zurich, Switzerland.
- Quirynen, R., Gros, S., Houska, B., & Diehl, M. (2016). Lifted collocation integrators for direct optimal control in ACADO toolkit. *Mathematical Programming Computation* (under review, preprint available at *Optimization Online* - www.optimization-online.org/DB_HTML/2016/05/5468.html).
- Quirynen, R., Vukov, M., & Diehl, M. (2012). Auto generation of implicit integrators for embedded NMPC with microsecond sampling times. In *Proceedings of the IFAC Nonlinear Model Predictive Control Conference* (pp. 175–180). Noordwijkerhout, The Netherlands.
- Quirynen, R., Vukov, M., Zanon, M., & Diehl, M. (2014). Auto-generating microsecond solvers for nonlinear MPC: A tutorial using ACADO integrators. *Optimal Control Applications and Methods*, 36, 685–704.
- Richter, S., Jones, C., & Morari, M. (2012). Computational complexity certification for real-time MPC with input constraints based on the fast gradient method. *IEEE Transactions on Automatic Control*, 57(6), 1391–1403.
- Santos, L. O., Afonso, P., Castro, J., Oliveira, N., & Biegler, L. T. (2001). On-line implementation of nonlinear MPC: An experimental case study. *Control Engineering Practice*, 9(8), 847–857.
- Shahzad, A., & Goulart, P. (2011). A new hot-start interior-point method for model predictive control. In *Proceedings of the 18th IFAC World Congress* (pp. 2470 – 2475). Milano, Italy.
- Shahzad, A., Kerrigan, E., & Constantinides, G. (2010). A warm-start interior-point method for predictive control. In *Proceedings of the UKACC International Conference on Control*. Coventry, UK.
- Siciliano, B., Sciavicco, L., Villani, L., & Oriolo, G. (2009). *Robotics: Modelling, planning and control*. Berlin: Springer.
- Simon, D., Loeferberg, J., & Glad, T. (2013). Nonlinear model predictive control using feedback linearization and local inner convex constraint approximations. In *Proceedings of the 12th European Control Conference* (pp. 2056–2061). Zurich, Switzerland.
- Vanderbei, R., & Shanno, D. (1999). An interior point algorithm for nonconvex nonlinear programming. *Computational Optimization and Applications*, 13, 231–252.
- Wächter, A., & Biegler, L. (2006). On the implementation of an interior-point filter line-search algorithm for large-scale nonlinear programming. *Mathematical Programming*, 106(25), 25–57.
- Wright, S. J. (1993). Interior point methods for optimal control of discrete time systems. *Journal of Optimization Theory and Applications*, 77(1), 161–187.
- Yildirim, E., & Wright, S. J. (2002). Warm-start strategies in interior-point methods for linear programming. *SIAM Journal on Optimization*, 12(3), 782–810.
- Zanelli, A., Quirynen, R., Jerez, J., & Diehl, M. (2017). A homotopy-based nonlinear interior-point method for NMPC applications. In *Proceedings of 20th IFAC World Congress*. Toulouse, France.
- Zavala, V. M., & Biegler, L. T. (2009). The advanced step NMPC controller: Optimality, stability and robustness. *Automatica*, 45(1), 86–93.