

MODEL PREDICTIVE CONTROL

DATA-DRIVEN MPC

Alberto Bemporad

`imt.lu/ab`

COURSE STRUCTURE

- ✓ Basic concepts of model predictive control (MPC) and linear MPC
- ✓ Linear time-varying and nonlinear MPC
- ✓ MPC computations: quadratic programming (QP), explicit MPC
- ✓ Hybrid MPC
- ✓ Stochastic MPC
- Data-driven MPC

Course page:


http://cse.lab.imtlucca.it/~bemporad/mpc_course.html

LEARNING MPC FROM DATA

- **Goal:** learn MPC law from data that optimizes a given index
- **Reinforcement learning** = use **data** and a **performance index** to learn an optimal policy
- **Q-learning:** learn Q-function defining the MPC law from data
(Gros, Zanon, 2019) (Zanon, Gros, Bemporad, 2019)
- **Policy gradient methods:** learn optimal policy coefficients directly from data using stochastic gradient descent (Ferrarotti, Bemporad, 2019)
- **Global optimization methods:** learn MPC parameters (weights, models, horizon, solver tolerances, ...) by optimizing observed closed-loop performance
(Piga, Forgione, Formentin, Bemporad, 2019) (Forgione, Piga, Bemporad, 2020)

DATA-DRIVEN MPC


DATA-DRIVEN MPC



- Can we design an MPC controller **without** first identifying a model of the **open-loop process**?

DATA-DRIVEN DIRECT CONTROLLER SYNTHESIS

(Campi, Lecchini, Savaresi, 2002) (Formentin et al., 2015)




- Collect a set of **data** $\{u(t), y(t), p(t)\}, t = 1, \dots, N$
- Specify a **desired closed-loop linear model** \mathcal{M} from r to y
- Compute $r_v(t) = \mathcal{M}^\# y(t)$ from **pseudo-inverse model** $\mathcal{M}^\#$ of \mathcal{M}
- **Identify** linear (LPV) model K_p from $e_v = r_v - y$ (virtual tracking error) to u

DATA-DRIVEN MPC

- Design a linear MPC (**reference governor**) to generate the reference r

(Bemporad, Mosca, 1994) (Gilbert, Kolmanovsky, Tan, 1994)





- MPC designed to handle input/output **constraints** and improve **performance**


(Piga, Formentin, Bemporad, 2017)

DATA-DRIVEN MPC - AN EXAMPLE

- Experimental results: MPC handles soft constraints on u , Δu and y
(motor equipment by courtesy of TU Delft)



desired tracking
performance achieved




constraints on input
increments satisfied

No open-loop process model is identified to design the MPC controller!

OPTIMAL DATA-DRIVEN MPC

- Question: How to choose the reference model \mathcal{M} ?



- Can we choose \mathcal{M} from data so that K_p is an **optimal controller**?

- **Idea:** parameterize desired closed-loop model $\mathcal{M}(\theta)$ and optimize

$$\min_{\theta} J(\theta) = \frac{1}{N} \sum_{t=0}^{N-1} \underbrace{W_y(r(t) - y_p(\theta, t))^2 + W_{\Delta u} \Delta u_p^2(\theta, t)}_{\text{performance index}} + \underbrace{W_{\text{fit}}(u(t) - u_v(\theta, t))^2}_{\text{identification error}}$$

- Evaluating $J(\theta)$ requires synthesizing $K_p(\theta)$ from data and simulating the nominal model and control law

$$y_p(\theta, t) = \mathcal{M}(\theta)r(t) \quad u_p(\theta, t) = K_p(\theta)(r(t) - y_p(\theta, t))$$


$$\Delta u_p(\theta, t) = u_p(\theta, t) - u_p(\theta, t-1)$$

- Optimal θ obtained by solving a **(non-convex) nonlinear programming** problem

- Results: **linear** process

$$G(z) = \frac{z - 0.4}{z^2 + 0.15z - 0.325}$$

Data-driven controller **only 1.3% worse** than model-based LQR (=SYS-ID on same data + LQR design)




- Results: **nonlinear (Wiener)** process

$$y_L(t) = G(z)u(t)$$

$$y(t) = |y_L(t)| \arctan(y_L(t))$$

The data-driven controller is **24% better** than LQR based on identified open-loop model !



DATA-DRIVEN OPTIMAL POLICY SEARCH

- Plant + environment dynamics (**unknown**):

$$s_{t+1} = h(s_t, p_t, u_t, d_t)$$

- s_t states of plant & environment
- p_t exogenous signal (e.g., reference)
- u_t control input
- d_t unmeasured disturbances

- Control policy**: $\pi : \mathbb{R}^{n_s+n_p} \longrightarrow \mathbb{R}^{n_u}$ deterministic control policy

$$u_t = \pi(s_t, p_t)$$

- Closed-loop **performance** of an execution is defined as

$$\mathcal{J}_{\infty}(\pi, s_0, \{p_{\ell}, d_{\ell}\}_{\ell=0}^{\infty}) = \sum_{\ell=0}^{\infty} \rho(s_{\ell}, p_{\ell}, \pi(s_{\ell}, p_{\ell}))$$

$$\rho(s_{\ell}, p_{\ell}, \pi(s_{\ell}, p_{\ell})) = \text{stage cost}$$

OPTIMAL POLICY SEARCH PROBLEM

- **Optimal policy:**

$$\begin{aligned}\pi^* &= \arg \min_{\pi} \mathcal{J}(\pi) \\ \mathcal{J}(\pi) &= \mathbb{E}_{s_0, \{p_\ell, d_\ell\}} [\mathcal{J}_\infty(\pi, s_0, \{p_\ell, d_\ell\})] \quad \text{expected performance}\end{aligned}$$

- **Simplifications:**

- Finite parameterization: $\pi = \pi_K(s_t, p_t)$ with K = parameters to optimize
- Finite horizon: $\mathcal{J}_L(\pi, s_0, \{p_\ell, d_\ell\}_{\ell=0}^{L-1}) = \sum_{\ell=0}^{L-1} \rho(s_\ell, p_\ell, \pi(s_\ell, p_\ell))$

- Optimal policy search: use **stochastic gradient descent (SGD)**


$$K_t \leftarrow K_{t-1} - \alpha_t \mathcal{D}(K_{t-1})$$

with $\mathcal{D}(K_{t-1})$ = descent direction

DESCENT DIRECTION

- The descent direction $\mathcal{D}(K_{t-1})$ is computed by generating:
 - N_s perturbations $s_0^{(i)}$ around the current state s_t
 - N_r random reference signals $r_\ell^{(j)}$ of length L ,
 - N_d random disturbance signals $d_\ell^{(h)}$ of length L ,

$$\mathcal{D}(K_{t-1}) = \sum_{i=1}^{N_s} \sum_{j=1}^{N_r} \sum_{k=1}^{N_d} \nabla_K \mathcal{J}_L(\pi_{K_{t-1}}, s_0^{(i)}, \{r_\ell^{(j)}, d_\ell^{(k)}\})$$



SGD step = mini-batch of size $M = N_s \cdot N_r \cdot N_d$

- Computing $\nabla_K \mathcal{J}_L$ requires predicting the effect of π over L future steps
- We use a **local linear model** just for computing $\nabla_K \mathcal{J}_L$, obtained by running **recursive linear system identification**

OPTIMAL POLICY SEARCH ALGORITHM

- At each step t :
 1. Acquire current s_t
 2. Recursively update the local linear model
 3. Estimate the direction of descent $\mathcal{D}(K_{t-1})$
 4. Update policy: $K_t \leftarrow K_{t-1} - \alpha_t \mathcal{D}(K_{t-1})$
- If policy is **learned online** and needs to be applied to the process:
 - Compute the nearest policy K_t^* to K_t that stabilizes the local model


$$K_t^* = \underset{K}{\operatorname{argmin}} \|K - K_t^s\|_2^2$$

s.t. K stabilizes local linear model *linear matrix inequality*

- When policy is learned online, **exploration** is guaranteed by the reference r_t

SPECIAL CASE: OUTPUT TRACKING

- $x_t = [y_t, y_{t-1}, \dots, y_{t-n_o}, u_{t-1}, u_{t-2}, \dots, u_{t-n_i}]$
 $\Delta u_t = u_t - u_{t-1}$ control input increment
- Stage cost: $\|y_{t+1} - r_t\|_{Q_y}^2 + \|\Delta u_t\|_R^2 + \|q_{t+1}\|_{Q_q}^2$
- Integral action dynamics $q_{t+1} = q_t + (y_{t+1} - r_t)$


$$s_t = \begin{bmatrix} x_t \\ q_t \end{bmatrix}, \quad p_t = r_t.$$

- **Linear policy parametrization:**


$$\pi_K(s_t, r_t) = -K^s \cdot s_t - K^r \cdot r_t, \quad K = \begin{bmatrix} K^s \\ K^r \end{bmatrix}$$

EXAMPLE: RETRIEVE LQR FROM DATA

$$\begin{cases} x_{t+1} &= \begin{bmatrix} -0.669 & 0.378 & 0.233 \\ -0.288 & -0.147 & -0.638 \\ -0.337 & 0.589 & 0.043 \end{bmatrix} x_t + \begin{bmatrix} -0.295 \\ -0.325 \\ -0.258 \end{bmatrix} u_t \\ y_t &= \begin{bmatrix} -1.139 & 0.319 & -0.571 \end{bmatrix} x_t \end{cases}$$

model is unknown

Online tracking performance (no disturbance, $d_t = 0$):




$$\begin{aligned} Q_y &= 1 \\ R &= 0.1 \\ Q_q &= 1 \end{aligned}$$

n_i	n_o	L
3	3	20
N_0	N_r	N_q
50	1	10

EXAMPLE: RETRIEVE LQR FROM DATA


Evolution of the error $\|K_t - K_{opt}\|_2$:



$$K_{\text{SGD}} = [-1.255, 0.218, 0.652, 0.895, 0.050, 1.115, -2.186]$$

$$K_{\text{opt}} = [-1.257, 0.219, 0.653, 0.898, 0.050, 1.141, -2.196]$$

NONLINEAR EXAMPLE



Continuously Stirred Tank Reactor (CSTR)^[1]

model is unknown

Feed:

- concentration: 10 kg mol/m^3
- temperature: 298.15 K

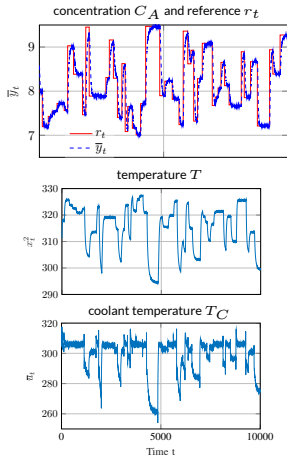
$$T = \hat{T} + \eta_T, \quad C_A = \hat{C}_A + \eta_C, \quad \eta_T, \eta_C \sim \mathcal{N}(0, \sigma^2), \quad \sigma = 0.01$$

$$Q_y = \begin{bmatrix} 1 & 0 \\ 0 & 0 \end{bmatrix} \quad R = 0.1 \quad Q_q = \begin{bmatrix} 0.01 & 0 \\ 0 & 0 \end{bmatrix}$$

[1] figure retrived from apmonitor.com

NONLINEAR EXAMPLE


Online learning





n_i	n_o	L
2	3	10
N_0	N_T	N_q
50	20	20

Validation phase

$$\text{Cost of } \mathbf{K}_{\text{SGD}} = 4.3 \cdot 10^3$$



$$\text{Cost of } \mathbf{K}_{\text{ID}} = 2.4 \cdot 10^4$$



Continuously Stirred Tank Reactor (CSTR)

(courtesy: apmonitor.com)

SGD beats SYS-ID + LQR

- Approach currently extended to multiple-linear and nonlinear policies


AUTOTUNING OF MPC

MPC CALIBRATION PROBLEM

- Controller depends on a vector x of parameters
- Parameters can be many things:
 - MPC weights, prediction model coefficients, horizons
 - Entries of covariance matrices in Kalman filter
 - Tolerances used in numerical solvers
 - ...
- Define a **performance index** f over a closed-loop simulation or real experiment.
For example:

$$f(x) = \sum_{t=0}^T \|y(t) - r(t)\|^2$$

(tracking quality)



- **Auto-tuning** = find the best combination of parameters that solves the **global optimization problem**

$$\min_x f(x)$$



What is a good optimization algorithm to solve $\min f(x)$?

- The algorithm should not require the gradient ∇f of $f(x)$
(**derivative-free** or **black-box optimization**)
- The algorithm should not get stuck on local minima (**global optimization**)
- The algorithm should make the **fewest evaluations** of the cost function f
(which is expensive to evaluate)

AUTO-TUNING - GLOBAL OPTIMIZATION ALGORITHMS


- Several derivative-free global optimization algorithms exist: (Rios, Sahidinis, 2013)
 - Lipschitzian-based partitioning techniques:
 - **DIRECT** (Divide in RECTangles) (Jones, 2001)
 - Multilevel Coordinate Search (**MCS**) (Huyer, Neumaier, 1999)
 - Response surface methods
 - **Kriging** (Matheron, 1967), **DACE** (Sacks et al., 1989)
 - Efficient global optimization (**EGO**) (Jones, Schonlau, Welch, 1998)
 - **Bayesian optimization** (Brochu, Cora, De Freitas, 2010)
 - Genetic algorithms (**GA**) (Holland, 1975)
 - Particle swarm optimization (**PSO**) (Kennedy, 2010)
 - ...
- **New method:** radial basis function surrogates + inverse distance weighting (**GLIS**) (Bemporad, 2019)

`cse.lab.imtlucca.it/~bemporad/glis`

- **Goal:** solve the **global optimization** problem

$$\begin{aligned} \min_x \quad & f(x) \\ \text{s.t.} \quad & \ell \leq x \leq u \\ & g(x) \leq 0 \end{aligned}$$

- **Step #0:** Get random initial samples $x_1, \dots, x_{N_{\text{init}}}$ (**Latin Hypercube Sampling**)



- **Step #1:** given N samples of f at x_1, \dots, x_N , build the **surrogate function**

$$\hat{f}(x) = \sum_{i=1}^N \beta_i \phi(\epsilon \|x - x_i\|_2)$$

ϕ = radial basis function

Example: $\phi(\epsilon d) = \frac{1}{1 + (\epsilon d)^2}$
(inverse quadratic)

Vector β solves $\hat{f}(x_i) = f(x_i)$ for all $i = 1, \dots, N$ (=linear system)

- **CAVEAT:** build and minimize $\hat{f}(x_i)$ iteratively may easily miss global optimum!


- **Step #2:** construct the **IDW exploration function**

$$z(x) = \frac{2}{\pi} \Delta F \tan^{-1} \left(\frac{1}{\sum_{i=1}^N w_i(x)} \right)$$

or 0 if $x \in \{x_1, \dots, x_N\}$

where $w_i(x) = \frac{e^{-\|x-x_i\|^2}}{\|x-x_i\|^2}$

ΔF = observed range of $f(x_i)$



- **Step #3:** optimize the **acquisition function**

$$x_{N+1} = \arg \min \quad \hat{f}(x) - \delta z(x)$$

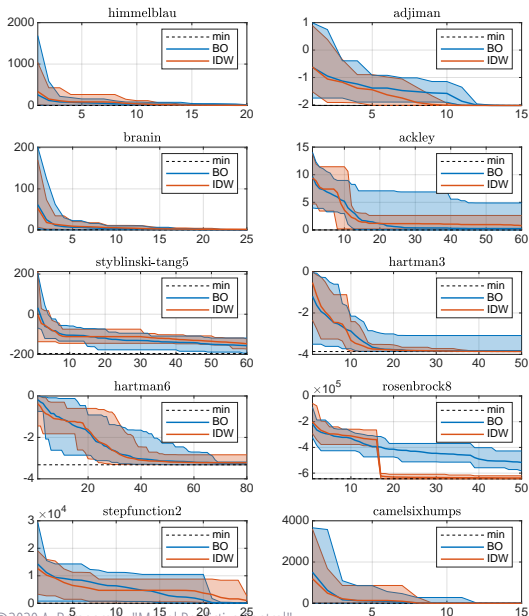
s.t. $\ell \leq x \leq u, g(x) \leq 0$

δ = exploitation vs
exploration tradeoff

to get new sample x_{N+1}

- Iterate the procedure to get new samples $x_{N+2}, \dots, x_{N_{\max}}$

GLIS VS BAYESIAN OPTIMIZATION



problem	n	BO[s]	IDW[s]
ackley	2	26.42	3.24
adjiman	2	3.39	0.66
branin	2	9.58	1.27
camelsixhumps	2	4.49	0.62
hartman3	3	23.19	3.58
hartman6	6	52.73	10.08
himmelblau	2	7.15	0.92
rosenbrock8	8	58.31	11.45
stepfunction2	4	10.52	1.72
styblinski-tang5	5	33.30	5.80


Results computed on 20 runs per test

BO = MATLAB's **bayesopt** fcn

AUTO-TUNING: MPC EXAMPLE

- We want to auto-tune the linear MPC controller




$$\begin{aligned}
 \min \quad & \sum_{k=0}^{50-1} (y_{k+1} - r(t))^2 + (W^{\Delta u} (u_k - u_{k-1}))^2 \\
 \text{s.t.} \quad & x_{k+1} = Ax_k + Bu_k \\
 & y_c = Cx_k \\
 & -1.5 \leq u_k \leq 1.5 \\
 & u_k \equiv u_{N_u}, \forall k = N_u, \dots, N-1
 \end{aligned}$$



- Calibration parameters: $x = [\log_{10} W^{\Delta u}, N_u]$
- Range: $-5 \leq x_1 \leq 3$ and $1 \leq x_2 \leq 50$
- Closed-loop performance objective:

$$f(x) = \sum_{t=0}^T \underbrace{(y(t) - r(t))^2}_{\text{track well}} + \underbrace{\frac{1}{2}(u(t) - u(t-1))^2}_{\text{smooth control action}} + \underbrace{2N_u}_{\text{small } QP}$$

AUTO-TUNING: EXAMPLE



• Result: $x^* = [-0.2341, 2.3007]$



$W^{\Delta u} = 0.5833, N_u = 2$

AUTO-TUNING: PROS AND CONS

- Pros:

- 👍 Selection of calibration parameters x to test is fully automatic
- 👍 Applicable to any calibration parameter (weights, horizons, solver tolerances, ...)
(Piga, Forgone, Formentin, Bemporad, 2019) (Forgione, Piga, Bemporad, 2020)
- 👍 Rather arbitrary performance index $f(x)$ (tracking performance, response time, worst-case number of flops, ...)

- Cons:

- 👎 Need to **quantify** an objective function $f(x)$
- 👎 No room for **qualitative** assessments of closed-loop performance
- 👎 Often objectives are multiple, not clear how to blend them in a **single** one

- Current research: **preference-based optimization (GLISp)**, having human assessments in the loop (**semi-automatic tuning**)

(Bemporad, Piga, 2019)

(Zhu, Bemporad, Piga, 2020)

cse.lab.imtlucca.it/~bemporad/glis

LEARNING MPC FROM DATA - LESSON LEARNED SO FAR

- Model/policy structure **includes** real plant/optimal policy:
 - **Sys-id + model-based** synthesis = model-free **reinforcement learning**
 - Reinforcement learning **may** require more data
(model-based can instead “extrapolate” optimal actions)
- Model/policy structure **does not include** real plant/optimal policy:
 - optimal policy **learned from data** **may** be better than **model-based** optimal policy
 - when open-loop model is used as a tuning parameter, **learned model** can be quite different from best **open-loop model** that can be identified from the same data